

Exploratory Data Analysis

Author: Raghu Raman Nanduri

Date: March 03,2019

Course: DSC530 - Term Project

Source Link: <https://www.kaggle.com/noriuk/us-education-datasets-unification-project/version/4> (<https://www.kaggle.com/noriuk/us-education-datasets-unification-project/version/4>)

Data set: finance_districts.csv

Statistical/Hypothetical Question: By exploring this data set regarding financial school districts and their enrollment numbers, I want to find out whether there is a any statistical correlation between total revenues of the state and number of enrollment numbers for the school districts. If there is correlation, how much have an impact Total Revenues have on enrollment.

Import libraries and data set

► In [11]:

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from math import sqrt
import scipy.stats as stats
import seaborn as sns

%matplotlib inline

plt.style.use('bmh')
```

```

In [12]: # Importing data set
findistdf = pd.read_csv('src/finance_districts.csv')
findistdf.head()

```

Out[12]:

	STATE	ENROLL	NAME	YRDATA	TOTALREV	TFEDREV	TSTREV	TLOCREV	TOTALEXP	TCURINST	TCURSSVC	TCURONON
0	ALABAMA	9609.0	AUTAUGA COUNTY SCHOOL DISTRICT	2016	80867	7447	53842	19578	76672	43843	23941	6401.0
1	ALABAMA	30931.0	BALDWIN COUNTY SCHOOL DISTRICT	2016	338236	23710	145180	169346	299880	164977	97231	19439.0
2	ALABAMA	912.0	BARBOUR COUNTY SCHOOL DISTRICT	2016	10116	2342	5434	2340	10070	4907	3896	975.0
3	ALABAMA	2842.0	EUFAULA CITY SCHOOL DISTRICT	2016	26182	3558	15900	6724	29843	15302	7901	2274.0
4	ALABAMA	3322.0	BIBB COUNTY SCHOOL DISTRICT	2016	32486	3664	21846	6976	31662	16407	11087	3122.0

Variables

```
► In [13]: findistdf.columns
```

```
Out[13]: Index(['STATE', 'ENROLL', 'NAME', 'YRDATA', 'TOTALREV', 'TFEDREV', 'TSTREV',  
              'TLOCREV', 'TOTALEXP', 'TCURINST', 'TCURSSVC', 'TCURONON', 'TCAPOUT'],  
              dtype='object')
```

Description of the data set

The data set consists of financials of each school district in each state for different years. It has the following variables:

STATE - State of Financial School district

ENROLL - The U.S. Census Bureau's count for students in the state. Should be comparable to GRADES_ALL

NAME - Name of the school district

YRDATA - Year that the record pertains to

TOTALREV: The total amount of revenue for the state.

 TFEDREV - Federal Revenue

 TSTREV - State Revenue

 TLOCREV - Local Revenue

TOTALEXP: The total expenditure for the state.

 TCURINST - Instruction Expenditure

 TCURSSVC - Supportive Services Expenditure

 TCURONON - Other Expenditure

 TCAPOUT - Capital Outlay Expenditure

Note: link for data set (finance_districts.csv)- <https://www.kaggle.com/adrian1acoran/starter-u-s-education-datasets-4a0c2b4b-7/data>

Data Cleansing

```
► In [14]: # Let us ignore the variables that are not part of this analysis
findistdf_orig = findistdf # taking backup of the original dataset

findistdf = findistdf[['STATE', 'ENROLL', 'NAME', 'YRDATA', 'TOTALREV', 'TFEDREV', 'TSTREV', 'TLOCREV', 'TOTAL
```

```
► In [15]: # Lets look at the basic information of the data set
findistdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 374161 entries, 0 to 374160
Data columns (total 9 columns):
STATE      374161 non-null object
ENROLL     358293 non-null float64
NAME       374161 non-null object
YRDATA     374161 non-null int64
TOTALREV   374161 non-null int64
TFEDREV    374161 non-null int64
TSTREV     374161 non-null int64
TLOCREV    374161 non-null int64
TOTALEXP   374161 non-null int64
dtypes: float64(1), int64(6), object(2)
memory usage: 25.7+ MB
```

Looks like ENROLL variable data is not available for the rows. For this analysis, discard the rows that have null or nan values in the 2 important variables - ENROLL and TOTALREV

```
► In [16]: findistdf = findistdf.dropna(subset = ['ENROLL', 'TOTALREV'])  
findistdf_orig2 = findistdf  
findistdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 358293 entries, 0 to 358292  
Data columns (total 9 columns):  
STATE      358293 non-null object  
ENROLL     358293 non-null float64  
NAME       358293 non-null object  
YRDATA     358293 non-null int64  
TOTALREV   358293 non-null int64  
TFEDREV    358293 non-null int64  
TSTREV     358293 non-null int64  
TLOCREV    358293 non-null int64  
TOTALEXP   358293 non-null int64  
dtypes: float64(1), int64(6), object(2)  
memory usage: 27.3+ MB
```

Distributions - Histograms, Outliers

► In [17]: *# Getting basic stats*
findistdf.describe()

Out[17]:

	ENROLL	YRDATA	TOTALREV	TFEDREV	TSTREV	TLOCREV	TOTALEXP
count	3.582930e+05	358293.000000	3.582930e+05	3.582930e+05	3.582930e+05	3.582930e+05	3.582930e+05
mean	3.134504e+03	2004.400390	3.173910e+04	2.689896e+03	1.473310e+04	1.431610e+04	3.209838e+04
std	1.402406e+04	6.849252	1.816932e+05	1.926575e+04	8.418726e+04	8.637972e+04	1.937165e+05
min	0.000000e+00	1993.000000	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
25%	2.760000e+02	1998.000000	3.299000e+03	1.730000e+02	1.354000e+03	1.125000e+03	3.231000e+03
50%	9.480000e+02	2004.000000	9.498000e+03	5.380000e+02	4.354000e+03	3.513000e+03	9.433000e+03
75%	2.607000e+03	2010.000000	2.575700e+04	1.677000e+03	1.156600e+04	1.101400e+04	2.585600e+04
max	1.077381e+06	2016.000000	2.744836e+07	3.120314e+06	1.056801e+07	1.514124e+07	2.962010e+07

Drawing Histograms and Density plots


In [18]:

```
plt.figure(figsize=(9, 8))  
#sns.distplot(findistdf['ENROLL'], color='b', bins=25, hist_kws={'alpha': 0.4});  
sns.distplot(findistdf['ENROLL'], color='b', hist = True, bins=25) #, hist_kws={'alpha': 0.4});
```

C:\Installed\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.


```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x2253014b5f8>



```
▶ In [19]: findistdf['ENROLL'].hist(bins = 80)  
#plt.Locator_params(nbins=20)  
plt.xlabel('Enroll')  
plt.ylabel('Frequency')  
plt.title('Histogram of ENROLL')
```

Out[19]: Text(0.5,1,'Histogram of ENROLL')





```
In [20]: plt.figure(figsize=(9, 8))
#sns.distplot(findistdf['ENROLL'], color='b', bins=25, hist_kws={'alpha': 0.4});
sns.distplot(findistdf['TOTALREV'], color='b', hist = True, bins=25) #, hist_kws={'alpha': 0.4});
```

C:\Installed\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.


```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x2252f574358>



```
► In [21]: # Trying to find apt number of bins from min to max
findistdf['TOTALREV'].hist(bins = 275, figsize=[12,6])
plt.locator_params(nbins=20)
plt.xlabel('Total Revenue')
plt.ylabel('Frequency')
plt.title('Histogram of TOTAL REVENUE')
```

Out[21]: Text(0.5,1,'Histogram of TOTAL REVENUE')



Both ENROLL and TOTALREV variables are asymmetrically (positively) skewed with longer tail towards the higher values.

Clearly there are some outliers in ENROLL and TOTALREV that are skewing the distributions heavily.

► In [22]: `findistdf.groupby('STATE').TOTALREV.agg(['min', 'max', 'mean', 'var'])`

Out[22]:

	min	max	mean	var
STATE				
ALABAMA	782	618181	4.349568e+04	4.077881e+09
ALASKA	158	914050	3.285153e+04	6.995753e+09
ARIZONA	0	635243	2.863350e+04	3.976721e+09
ARKANSAS	0	371708	1.339582e+04	6.637229e+08
CALIFORNIA	0	10329380	5.464825e+04	6.051548e+10
COLORADO	0	1271873	3.538293e+04	1.027157e+10
CONNECTICUT	78	527210	4.415065e+04	3.689717e+09
DELAWARE	3879	318466	7.314053e+04	4.044500e+09
DISTRICT_OF_COLUMBIA	678874	1382282	1.033762e+06	6.165597e+10
FLORIDA	5712	3959408	3.172400e+05	3.023165e+11
GEORGIA	739	1893150	7.196012e+04	2.694898e+10
HAWAII	1062475	3030519	2.047635e+06	4.478832e+11
IDAHO	41	281989	1.491776e+04	8.556664e+08
ILLINOIS	0	5760419	2.163508e+04	1.942673e+10
INDIANA	0	729047	3.179480e+04	2.305822e+09
IOWA	683	482554	1.250332e+04	6.858602e+08
KANSAS	551	666255	1.472987e+04	1.616821e+09
KENTUCKY	0	1323404	3.216628e+04	6.154377e+09
LOUISIANA	7144	625708	9.061099e+04	1.280243e+10
MAINE	0	119828	7.764532e+03	1.349948e+08
MARYLAND	16763	3159510	4.132408e+05	3.104588e+11
MASSACHUSETTS	0	1481699	3.228844e+04	4.637021e+09
MICHIGAN	0	1671978	2.475151e+04	3.685789e+09

	min	max	mean	var
STATE				
MINNESOTA	0	716799	2.087318e+04	2.691517e+09
MISSISSIPPI	0	291507	2.321156e+04	8.352367e+08
MISSOURI	0	510063	1.506467e+04	1.466022e+09
MONTANA	0	117079	2.835495e+03	4.497919e+07
NEBRASKA	1	687093	6.218597e+03	8.453300e+08
NEVADA	1294	3220684	1.803720e+05	2.960478e+11
NEW_HAMPSHIRE	2	186461	1.244669e+04	3.501079e+08
NEW_JERSEY	1	1163400	3.597258e+04	4.114898e+09
NEW_MEXICO	745	1051179	3.249878e+04	7.939642e+09
NEW_YORK	1	27448356	6.258730e+04	4.314818e+11
NORTH_CAROLINA	4938	2297008	9.167576e+04	2.673796e+10
NORTH_DAKOTA	10	177525	4.296540e+03	1.610324e+08
OHIO	1	980035	2.541799e+04	2.812349e+09
OKLAHOMA	0	478115	8.822205e+03	6.638011e+08
OREGON	39	697501	2.195237e+04	2.554278e+09
PENNSYLVANIA	0	3030964	3.618602e+04	9.705357e+09
RHODE_ISLAND	1732	461090	4.848231e+04	3.628086e+09
SOUTH_CAROLINA	0	855902	6.281708e+04	9.269236e+09
SOUTH_DAKOTA	17	230640	6.363154e+03	2.463003e+08
TENNESSEE	0	1470379	4.814285e+04	1.225541e+10
TEXAS	0	2480131	3.575125e+04	1.146961e+10
UTAH	1812	635282	8.064527e+04	1.400811e+10
VERMONT	0	84611	4.518641e+03	3.780746e+07
VIRGINIA	0	2733933	8.174012e+04	4.034570e+10
WASHINGTON	25	881789	3.143074e+04	3.719914e+09
WEST_VIRGINIA	2522	384864	4.753317e+04	2.279495e+09

	min	max	mean	var
STATE				
WISCONSIN	493	1310838	2.122801e+04	3.447978e+09
WYOMING	593	309591	2.482201e+04	1.315037e+09

```

In [23]: findist_bystate = pd.DataFrame()
findist_bystate = findistdf.groupby('STATE', as_index = False)['TOTALREV', 'ENROLL'].max()

```

```

In [24]: #Top 5 states with most total revenue
findist_bystate.sort_values('TOTALREV',ascending=False)[0:5]

```

Out[24]:

	STATE	TOTALREV	ENROLL
32	NEW_YORK	27448356	1077381.0
4	CALIFORNIA	10329380	747009.0
13	ILLINOIS	5760419	437418.0
9	FLORIDA	3959408	375836.0
28	NEVADA	3220684	325990.0

```

In [25]: #Top 5 states with most total revenue
findist_bystate.sort_values('ENROLL',ascending=False)[0:5]

```

Out[25]:

	STATE	TOTALREV	ENROLL
32	NEW_YORK	27448356	1077381.0
4	CALIFORNIA	10329380	747009.0
13	ILLINOIS	5760419	437418.0
9	FLORIDA	3959408	375836.0
28	NEVADA	3220684	325990.0

Clearly California and NewYork are in a different league with respect to Total Revenue and Enroll numbers. So, let's see how the distribution will be if we separate these states from the data set

```

In [26]: findistdf_orig3 = findistdf # taking backup copy of df

# separating the records of 'New York' and 'California' from the data set
findistdf = findistdf[(findistdf.STATE != 'NEW_YORK') & (findistdf.STATE != 'CALIFORNIA' )]

# removing recrods with 0 revenue's and enrollments

findistdf = findistdf[(findistdf.ENROLL >0) & (findistdf.TOTALREV >0) & (findistdf.TFEDREV >0)
                      & (findistdf.TSTREV >0) & (findistdf.TLOCREV >0) & (findistdf.TOTALEXP >0) ]

len(findistdf)

```

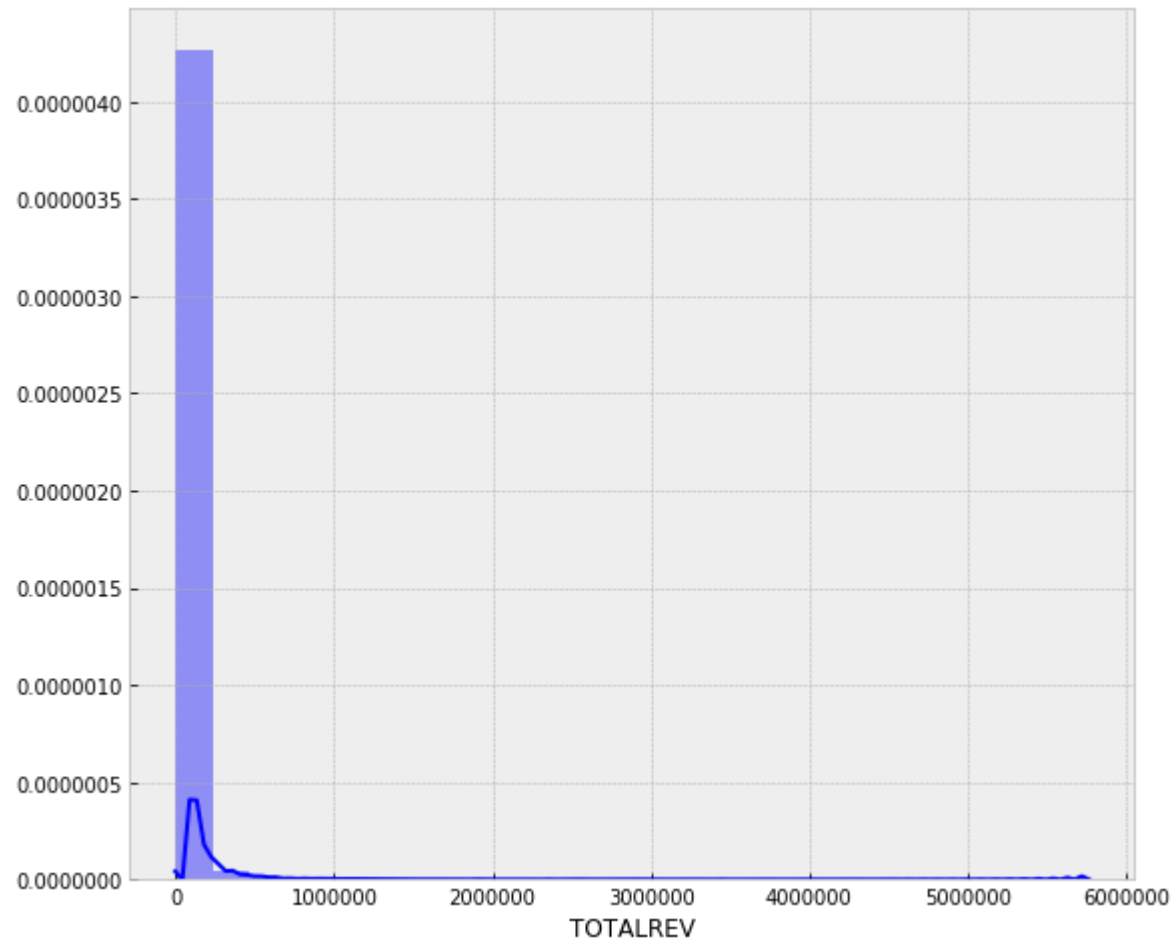
Out[26]: 287744

Plotting Density plots, histograms

```
► In [27]: # Density plot for TOTALREV
plt.figure(figsize=(9, 8))
sns.distplot(findistdf['TOTALREV'], color='b', bins=25, hist_kws={'alpha': 0.4});
```

C:\Installed\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```




```
► In [28]: # Trying to find apt number of bins from min to max
noofbins = np.arange(start=findistdf['TOTALREV'].min(), stop=findistdf['TOTALREV'].max(), step=50000)
print(len(noofbins))
```

116

```
► In [29]: findistdf['TOTALREV'].hist(bins = 25, figsize=[12,6])
plt.locator_params(nbins=20)
plt.xlabel('Total Revenue')
plt.ylabel('Frequency')
plt.title('Histogram of TOTAL REVENUE')
```


Out[29]: Text(0.5,1,'Histogram of TOTAL REVENUE')




```
► In [30]: # Density plot for ENROLL
plt.figure(figsize=(9, 8))
sns.distplot(findistdf['ENROLL'], color='b', bins=25, hist_kws={'alpha': 0.4});
```

C:\Installed\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

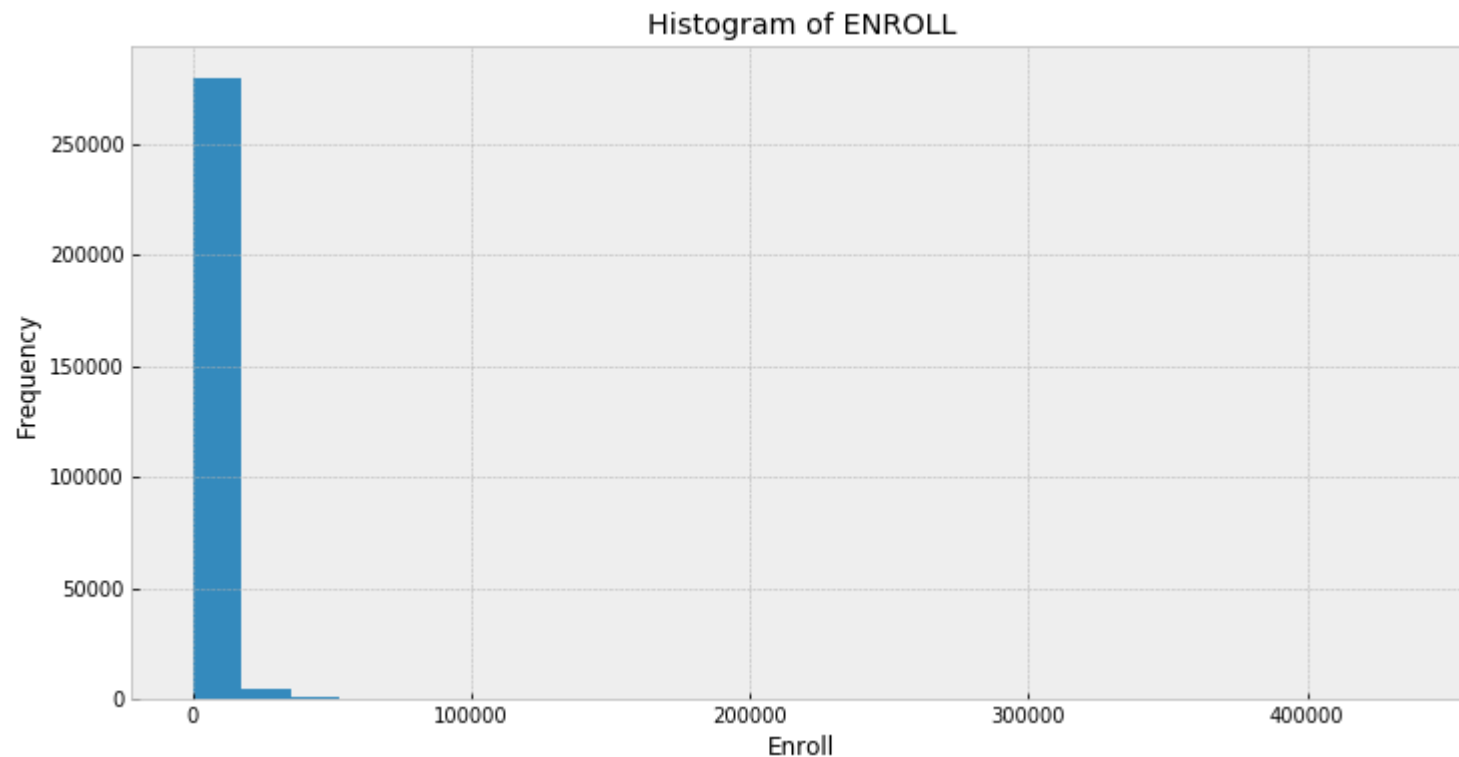


```
► In [31]: # Trying to find apt number of bins from min to max
noofbins = np.arange(start=findistdf['ENROLL'].min(), stop=findistdf['ENROLL'].max(), step=10000)
print(len(noofbins))
```

44

```
► In [32]: findistdf['ENROLL'].hist(bins = 25, figsize=[12,6])
#plt.locator_params(nbins=20)
plt.xlabel('Enroll')
plt.ylabel('Frequency')
plt.title('Histogram of ENROLL')
```

Out[32]: Text(0.5,1,'Histogram of ENROLL')



Still most of the numbers are packed at the lower end of the scale for the TOTALREV and ENROLL

Plotting Histogram for the 5 variables


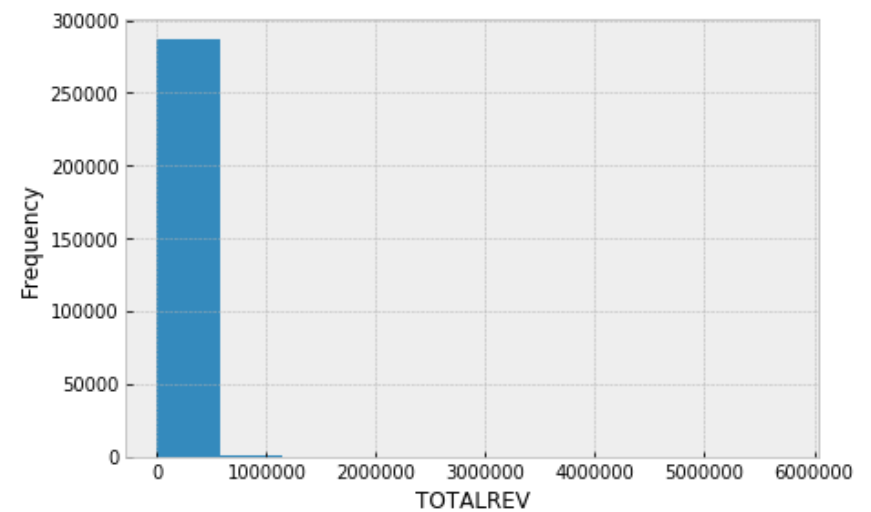
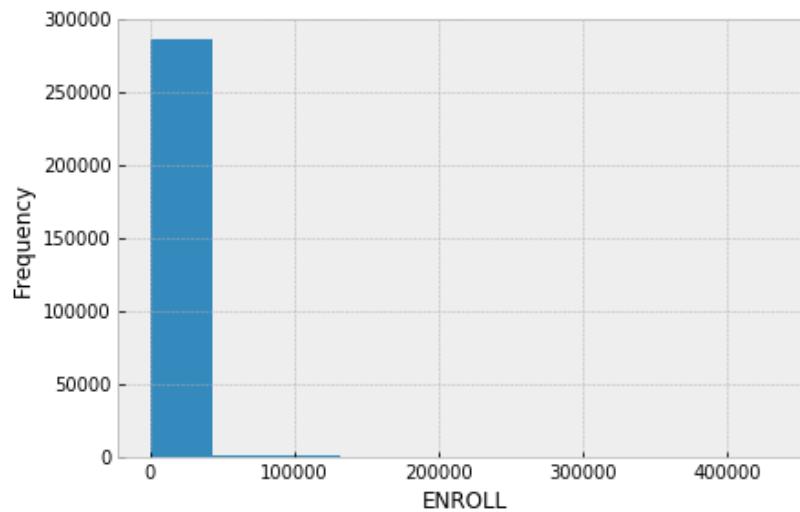
```

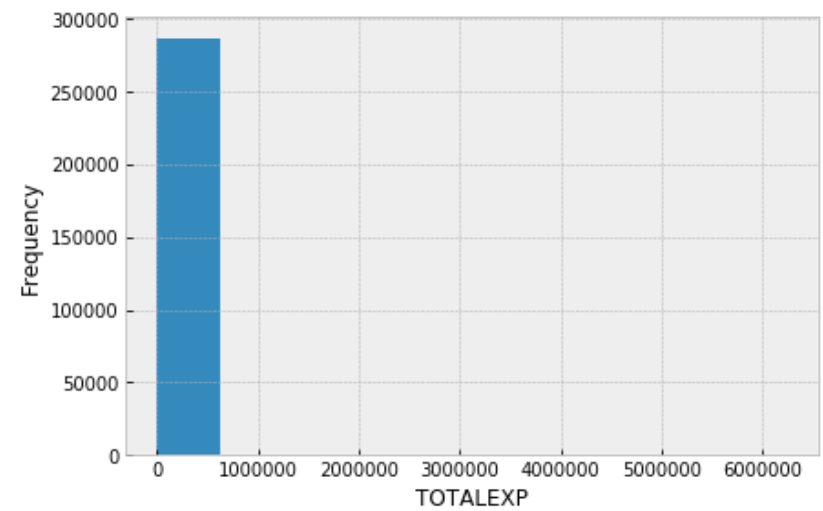

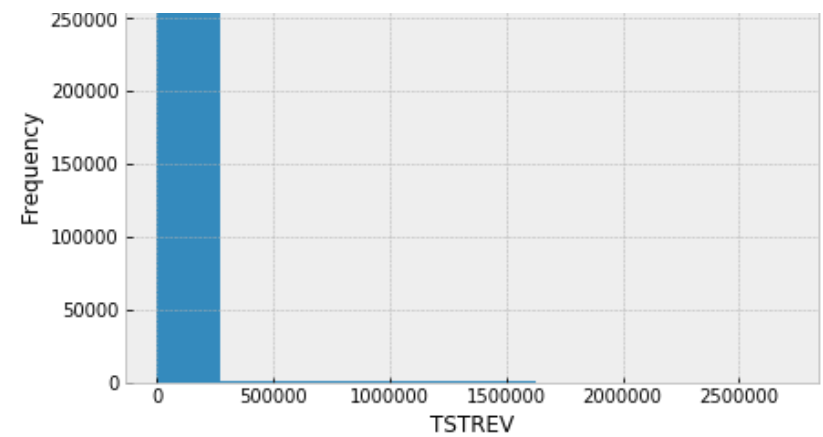
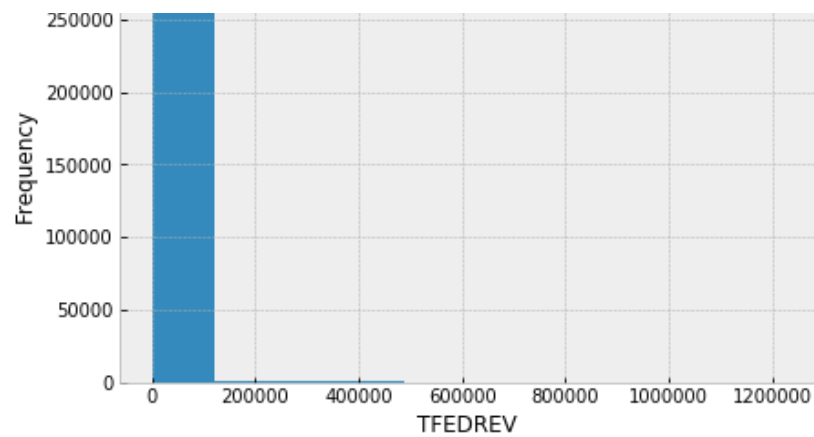
In [33]: colnames = ['ENROLL', 'TOTALREV', 'TFEDREV', 'TSTREV', 'TLOCREV', 'TOTALEXP']
plt.figure(figsize=(15,15))
for i in range(len(colnames)):

    if i == 0:
        plt.subplot(3,2,1)
    if i == 1:
        plt.subplot(3,2,2)
    if i == 2:
        plt.subplot(3,2,3)
    if i == 3:
        plt.subplot(3,2,4)
    if i == 4:
        plt.subplot(3,2,5)
    if i == 5:
        plt.subplot(3,2,6)

    #print(colnames[i])
    plt.hist(findistdf[colnames[i]].dropna())
    plt.ylabel('Frequency')
    plt.xlabel(findistdf[colnames[i]].name)

```





► In []:

Plotting Desnity plots for the 5 variables

```

In [34]: colnames = ['ENROLL', 'TOTALREV', 'TFEDREV', 'TSTREV', 'TLOCREV', 'TOTALEXP']
plt.figure(figsize=(15,15))
for i in range(len(colnames)):



    if i == 0:
        plt.subplot(3,2,1)
    if i == 1:
        plt.subplot(3,2,2)
    if i == 2:
        plt.subplot(3,2,3)
    if i == 3:
        plt.subplot(3,2,4)
    if i == 4:
        plt.subplot(3,2,5)
    if i == 5:
        plt.subplot(3,2,6)

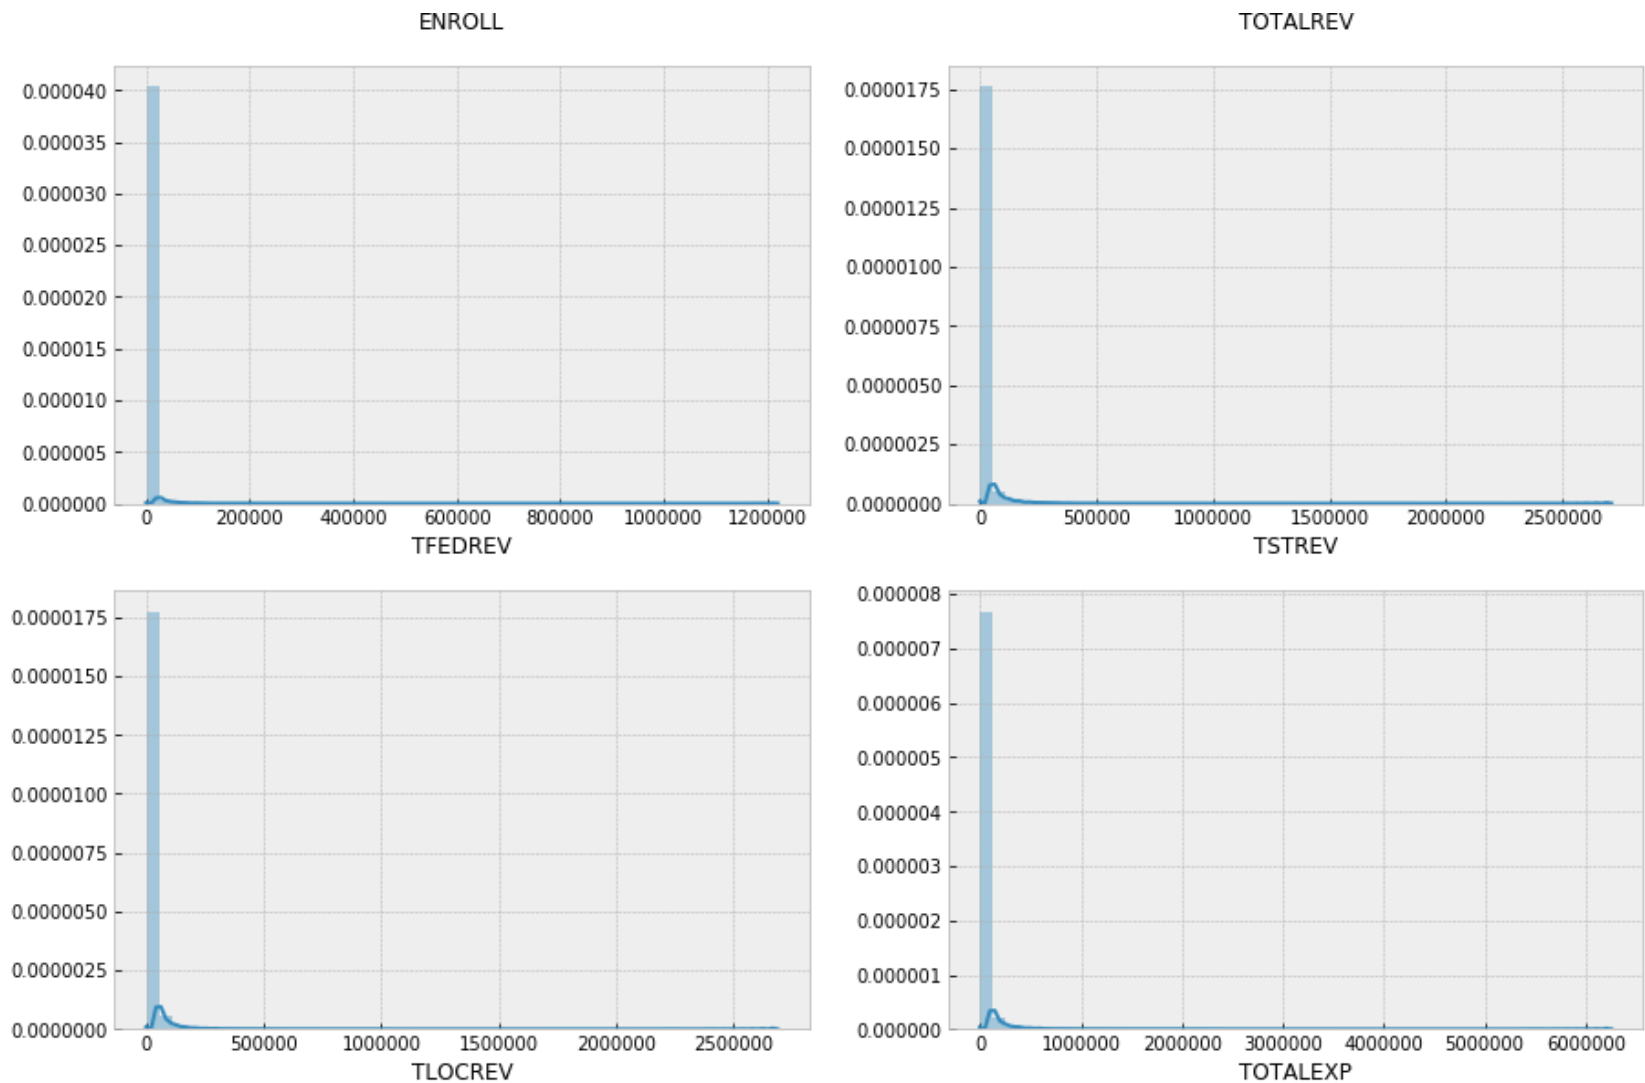
    #print(colnames[i])
    sns.distplot(findistdf[colnames[i]].dropna())

```

C:\Installed\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```





Plotting outliers using boxplot



```


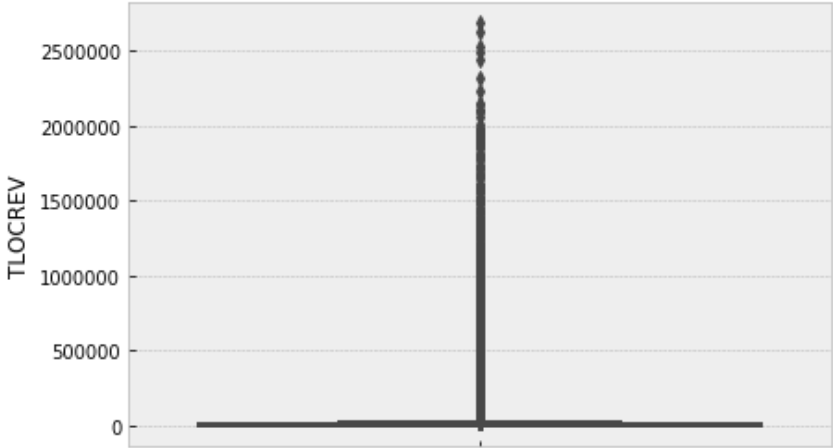

In [35]: plt.figure(figsize=(15,15))
for i in range(len(colnames)):

    if i == 0:
        plt.subplot(3,2,1)
    if i == 1:
        plt.subplot(3,2,2)
    if i == 2:
        plt.subplot(3,2,3)
    if i == 3:
        plt.subplot(3,2,4)
    if i == 4:
        plt.subplot(3,2,5)
    if i == 5:
        plt.subplot(3,2,6)

    #print(colnames[i])
    sns.boxplot(y=findistdf[colnames[i]].dropna())
    #column=findistdf[colnames[i]]
    #column.to_frame().boxplot(figsize=[4,8])
    #findistdf.boxplot(column=colnames[i], figsize=[4,8]);

```





Clearly the distribution of the metrics is not normal and they are absolutely skewed towards with the lower end of the scale with long tail on right end of the scale.

Let's transform the metrics into log form and see, how their distributions and histograms look like

```

In [36]: # Adding log transformed columns to the dataframe
findistdf['lg_ENROLL'] = findistdf['ENROLL'].apply(np.log)
findistdf['lg_TOTALREV'] = findistdf['TOTALREV'].apply(np.log)
findistdf['lg_TFEDREV'] = findistdf['TFEDREV'].apply(np.log)
findistdf['lg_TSTREV'] = findistdf['TSTREV'].apply(np.log)
findistdf['lg_TLOCREV'] = findistdf['TLOCREV'].apply(np.log)
findistdf['lg_TOTALEXP'] = findistdf['TOTALEXP'].apply(np.log)
findistdf.head()

```

Out[36]:

	STATE	ENROLL	NAME	YRDATA	TOTALREV	TFEDREV	TSTREV	TLOCREV	TOTALEXP	lg_ENROLL	lg_TOTALREV	lg_TFED
0	ALABAMA	9609.0	AUTAUGA COUNTY SCHOOL DISTRICT	2016	80867	7447	53842	19578	76672	9.170455	11.300561	8.911
1	ALABAMA	30931.0	BALDWIN COUNTY SCHOOL DISTRICT	2016	338236	23710	145180	169346	299880	10.339514	12.731499	10.071
2	ALABAMA	912.0	BARBOUR COUNTY SCHOOL DISTRICT	2016	10116	2342	5434	2340	10070	6.815640	9.221874	7.751
3	ALABAMA	2842.0	EUFAULA CITY SCHOOL DISTRICT	2016	26182	3558	15900	6724	29843	7.952263	10.172827	8.171
4	ALABAMA	3322.0	BIBB COUNTY SCHOOL DISTRICT	2016	32486	3664	21846	6976	31662	8.108322	10.388565	8.201

Plotting histograms of log transformed columns



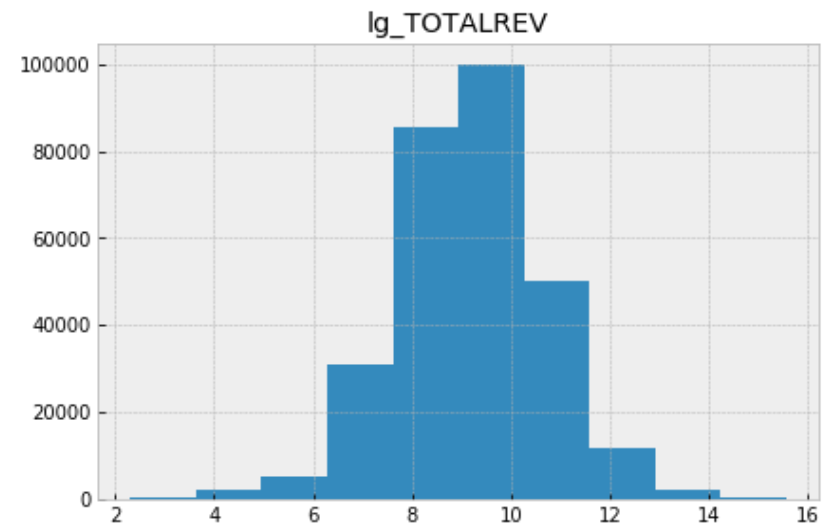
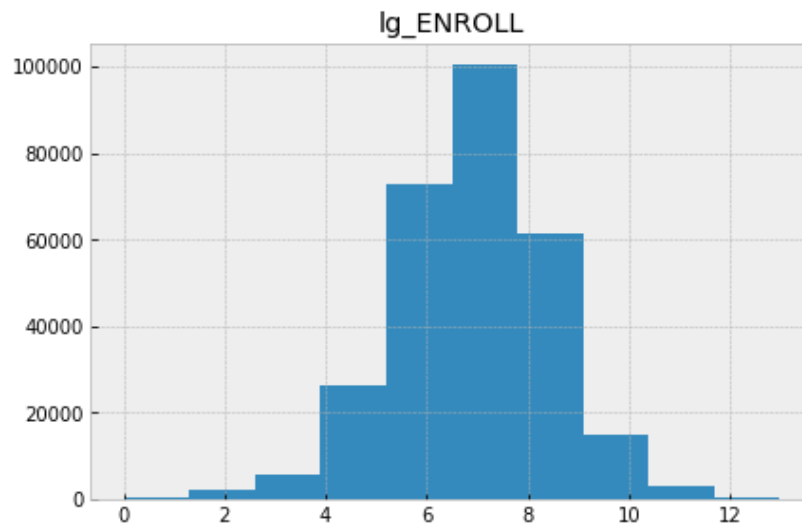
```

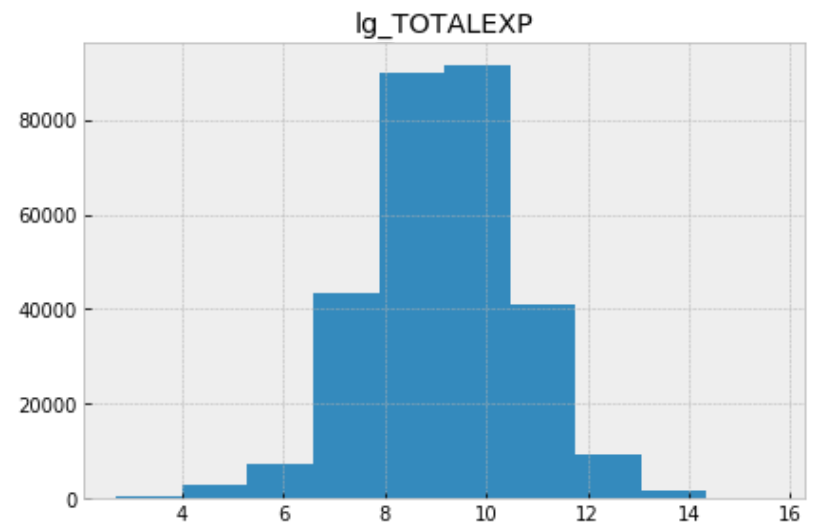



In [37]: lgcolnames = ['lg_ENROLL', 'lg_TOTALREV', 'lg_TFEDREV', 'lg_TSTREV', 'lg_TLOCREV', 'lg_TOTALEXP']
plt.figure(figsize=(15,15))
for i in range(len(lgcolnames)):

    if i == 0:
        plt.subplot(3,2,1)
    if i == 1:
        plt.subplot(3,2,2)
    if i == 2:
        plt.subplot(3,2,3)
    if i == 3:
        plt.subplot(3,2,4)
    if i == 4:
        plt.subplot(3,2,5)
    if i == 5:
        plt.subplot(3,2,6)

    #print(colnames[i])
    plt.hist(findistdf[lgcolnames[i]].astype('float'))
    plt.title(findistdf[lgcolnames[i]].name)

```





Histograms of log transformed variables doesn't appear as much skewed as they were earlier without log transformations. All variables except for lg_TOTALEXP appear in unimodal distribution where as lg_TOTALEXP is in bimodal distribution.



```
In [38]: plt.figure(figsize=(15,15))
for i in range(len(lgcolnames)):



    if i == 0:
        plt.subplot(3,2,1)
    if i == 1:
        plt.subplot(3,2,2)
    if i == 2:
        plt.subplot(3,2,3)
    if i == 3:
        plt.subplot(3,2,4)
    if i == 4:
        plt.subplot(3,2,5)
    if i == 5:
        plt.subplot(3,2,6)

    #print(colnames[i])
    sns.distplot(findistdf[lgcolnames[i]].dropna().astype('float'))
    plt.title(findistdf[lgcolnames[i]].name)
```

C:\Installed\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```











As per the above density plots, all the variable appear close to log normal distribution

```
▶ In [39]: plt.figure(figsize=(15,15))
for i in range(len(lgcolnames)):

    if i == 0:
        plt.subplot(3,2,1)
    if i == 1:
        plt.subplot(3,2,2)
    if i == 2:
        plt.subplot(3,2,3)
    if i == 3:
        plt.subplot(3,2,4)
    if i == 4:
        plt.subplot(3,2,5)
    if i == 5:
        plt.subplot(3,2,6)

    column=findistdf[lgcolnames[i]]
    column.to_frame().boxplot(figsize=[4,8])
```



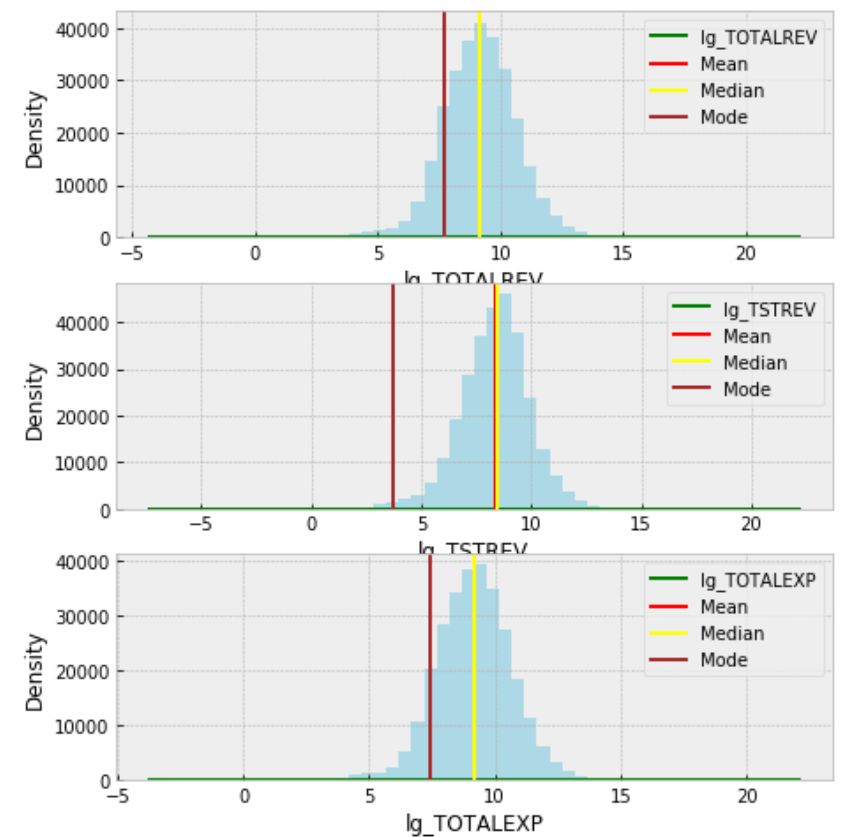
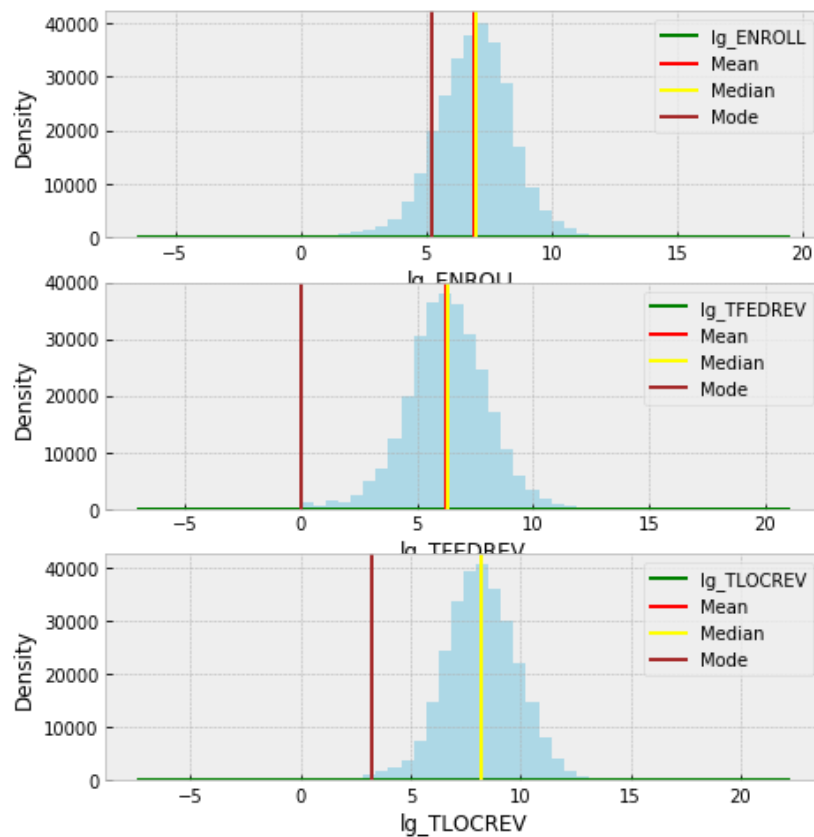


Log transformation gives a better representation of these variables with lesser outliers


```
► In [40]: #Plotting histogram and KDE for total revenue
plt.figure(figsize=(30, 30))
for i in range(len(lgcolnames)):

    if i == 0:
        plt.subplot(3,2,1)
    if i == 1:
        plt.subplot(3,2,2)
    if i == 2:
        plt.subplot(3,2,3)
    if i == 3:
        plt.subplot(3,2,4)
    if i == 4:
        plt.subplot(3,2,5)
    if i == 5:
        plt.subplot(3,2,6)

    axtr = findistdf[lgcolnames[i]].astype('float').hist(bins = 26, color = 'lightblue') #, normed=True)
    findistdf[lgcolnames[i]].plot(kind='kde', color='Green', ax=axtr, figsize=[16,8])
    #plt.locator_params(nbins=20)
    #plt.title('Histogram - KDE for %s with mean(red), median(yellow) and mode(brown), %findistdf[lgcolnames[i]
    plt.xlabel(findistdf[lgcolnames[i]].name);
    plt.axvline(findistdf[lgcolnames[i]].mean(),color='red',label='Mean')
    plt.axvline(findistdf[lgcolnames[i]].median(),color='yellow',label='Median')
    plt.axvline(findistdf[lgcolnames[i]].mode()[0],color='brown',label='Mode')
    plt.legend()
    #plt.title('Histogram - KDE with mean(red), median(yellow) and mode(brown)')
```



► In []:



Exploring data set by STATE and YEAR

► In [41]:

```
findist_bystateyr = pd.DataFrame()
findist_bystateyr = findistdf.groupby(['YRDATA', 'STATE']).sum()
```

► In [42]: *# Displaying top 5 total revenues by state and year*
 findist_bystateyr.sort_values('TOTALREV',ascending=False)[0:5]

Out[42]:

		ENROLL	TOTALREV	TFEDREV	TSTREV	TLOCREV	TOTALEXP	lg_ENROLL	lg_TOTALREV	lg_TFEDREV	lg_TSTR
YRDATA	STATE										
2016	TEXAS	5053260.0	57724562	5666460	21937670	30120432	58859458	7242.432264	9874.110413	7159.541881	8807.2995
2015	TEXAS	5004839.0	55046828	5550680	20903361	28592787	55822147	7232.728807	9847.412237	7121.440678	8722.9515
2014	TEXAS	4949437.0	52271215	5389557	20452946	26428712	52651972	7227.241208	9813.737375	7108.461573	8738.9475
2011	TEXAS	4800196.0	50448814	7542016	19510784	23396014	52221814	7239.639505	9827.987325	7568.709826	8848.1646
2010	TEXAS	4728815.0	49998660	7710320	18799593	23488747	53344838	7238.895148	9833.663269	7655.391596	8829.5895

► In [43]: *# Displaying bottom 5 total revenues by state and year*
 findist_bystateyr.sort_values('TOTALREV')[0:5]

Out[43]:

		ENROLL	TOTALREV	TFEDREV	TSTREV	TLOCREV	TOTALEXP	lg_ENROLL	lg_TOTALREV	lg_TFEDREV	
YRDATA	STATE										
1993	SOUTH_DAKOTA	107050.0	465650	52787	120851	292012	481665	535.762896	664.100837	471.248246	
1994	SOUTH_DAKOTA	108439.0	507927	51846	126591	329490	525289	536.448805	668.787000	470.863467	
1993	NORTH_DAKOTA	118536.0	549529	52163	231963	265403	557977	1281.962611	1705.264803	1046.491984	1
1994	NORTH_DAKOTA	118670.0	560962	50679	234002	276281	578406	1257.753819	1676.509178	1012.583536	1
1995	NORTH_DAKOTA	119132.0	600172	57770	244283	298119	580340	1214.217840	1623.333995	999.034394	1

In []:

Calculating Mean, Mode, Spread, and Tails

```
In [44]: # statistics for whole data set(including California & New York - school districts)
findistdf_orig2.describe()
```

Out[44]:

	ENROLL	YRDATA	TOTALREV	TFEDREV	TSTREV	TLOCREV	TOTALEXP
count	3.582930e+05	358293.000000	3.582930e+05	3.582930e+05	3.582930e+05	3.582930e+05	3.582930e+05
mean	3.134504e+03	2004.400390	3.173910e+04	2.689896e+03	1.473310e+04	1.431610e+04	3.209838e+04
std	1.402406e+04	6.849252	1.816932e+05	1.926575e+04	8.418726e+04	8.637972e+04	1.937165e+05
min	0.000000e+00	1993.000000	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
25%	2.760000e+02	1998.000000	3.299000e+03	1.730000e+02	1.354000e+03	1.125000e+03	3.231000e+03
50%	9.480000e+02	2004.000000	9.498000e+03	5.380000e+02	4.354000e+03	3.513000e+03	9.433000e+03
75%	2.607000e+03	2010.000000	2.575700e+04	1.677000e+03	1.156600e+04	1.101400e+04	2.585600e+04
max	1.077381e+06	2016.000000	2.744836e+07	3.120314e+06	1.056801e+07	1.514124e+07	2.962010e+07

```
▶ In [45]: def descchar(var):  
    """  
    Function to print descriptive characteristics of the variable  
    Input: dataframe.columnname  
    Returns: descriptive characteristics like - mean, median, mode, spread, interquartile range, skew  
    """  
  
    print(' Mean, Median, Mode of %s, %f %f %f ' %(var.name, var.mean(), var.median(), var.mode()[0]))  
    print(' Spread - Variance, Standard deviation of %s, %f %f ' %(var.name, var.var(), var.std()))  
    print(' Skew of %s, %f ' %(var.name, var.skew()))  
    print(' Interquartile range of %s, %f %f %f' %(var.name, var.quantile(0.25), var.quantile(0.5), var.quantile(0.75)))
```

```
► In [46]: # statistics for entire data set with California & New York school districts
for i in range(len(colnames)):
    print("Descriptive Characteristics for %s" % findistdf_orig2[colnames[i]].name)
    descchar(findistdf_orig2[colnames[i]])
```

Descriptive Characteristics for ENROLL

Mean, Median, Mode of ENROLL, 3134.504099 948.000000 0.000000

Spread - Variance, Standard deviation of ENROLL, 196674167.451108 14024.056740

Skew of ENROLL, 39.215369

Interquartile range of ENROLL, 276.000000 948.000000 2607.000000

Descriptive Characteristics for TOTALREV

Mean, Median, Mode of TOTALREV, 31739.097055 9498.000000 0.000000

Spread - Variance, Standard deviation of TOTALREV, 33012411727.751137 181693.180191

Skew of TOTALREV, 75.790358

Interquartile range of TOTALREV, 3299.000000 9498.000000 25757.000000

Descriptive Characteristics for TFEDREV

Mean, Median, Mode of TFEDREV, 2689.895753 538.000000 0.000000

Spread - Variance, Standard deviation of TFEDREV, 371169206.807162 19265.752173

Skew of TFEDREV, 62.346503

Interquartile range of TFEDREV, 173.000000 538.000000 1677.000000

Descriptive Characteristics for TSTREV

Mean, Median, Mode of TSTREV, 14733.104367 4354.000000 0.000000

Spread - Variance, Standard deviation of TSTREV, 7087494107.262826 84187.256205

Skew of TSTREV, 64.036625

Interquartile range of TSTREV, 1354.000000 4354.000000 11566.000000

Descriptive Characteristics for TLOCREV

Mean, Median, Mode of TLOCREV, 14316.096912 3513.000000 0.000000

Spread - Variance, Standard deviation of TLOCREV, 7461456055.094042 86379.720161

Skew of TLOCREV, 88.592093

Interquartile range of TLOCREV, 1125.000000 3513.000000 11014.000000

Descriptive Characteristics for TOTALEXP

Mean, Median, Mode of TOTALEXP, 32098.381944 9433.000000 0.000000

Spread - Variance, Standard deviation of TOTALEXP, 37526083710.425728 193716.503454

Skew of TOTALEXP, 79.133328

Interquartile range of TOTALEXP, 3231.000000 9433.000000 25856.000000

► In [47]: *# statistics for data set without California & New York school districts*
 findistdf.describe()

Out[47]:

	ENROLL	YRDATA	TOTALREV	TFEDREV	TSTREV	TLOCREV	TOTALEXP	Ig_ENROLL	Ig_TOTALRE
count	287744.000000	287744.000000	2.877440e+05	2.877440e+05	2.877440e+05	2.877440e+05	2.877440e+05	287744.000000	287744.000000
mean	3170.830384	2004.461594	3.025546e+04	2.472747e+03	1.398188e+04	1.380084e+04	3.056100e+04	6.912037	9.17085
std	10435.217797	6.848307	1.033309e+05	1.225883e+04	4.819758e+04	5.112197e+04	1.057474e+05	1.520550	1.4830
min	1.000000	1993.000000	1.000000e+01	1.000000e+00	1.000000e+00	1.000000e+00	1.500000e+01	0.000000	2.3025
25%	388.000000	1999.000000	3.653000e+03	1.870000e+02	1.623000e+03	1.224000e+03	3.592000e+03	5.961005	8.2033
50%	1073.000000	2004.000000	9.715000e+03	5.390000e+02	4.608000e+03	3.524000e+03	9.660000e+03	6.978214	9.1814
75%	2720.000000	2010.000000	2.520025e+04	1.649000e+03	1.150125e+04	1.072525e+04	2.534300e+04	7.908387	10.1346
max	437418.000000	2016.000000	5.760419e+06	1.220298e+06	2.710361e+06	2.687925e+06	6.253045e+06	12.988645	15.5665


```

In [48]: # statistics for data set without California & New York school districts
for i in range(len(colnames)):
    print("Descriptive Characteristics for %s" % findistdf[colnames[i]].name)
    descchar(findistdf[colnames[i]])

```

Descriptive Characteristics for ENROLL

Mean, Median, Mode of ENROLL, 3170.830384 1073.000000 180.000000
 Spread - Variance, Standard deviation of ENROLL, 108893770.472233 10435.217797
 Skew of ENROLL, 16.793895
 Interquartile range of ENROLL, 388.000000 1073.000000 2720.000000

Descriptive Characteristics for TOTALREV

Mean, Median, Mode of TOTALREV, 30255.458977 9715.000000 2276.000000
 Spread - Variance, Standard deviation of TOTALREV, 10677274625.544109 103330.898697
 Skew of TOTALREV, 18.316904
 Interquartile range of TOTALREV, 3653.000000 9715.000000 25200.250000

Descriptive Characteristics for TFEDREV

Mean, Median, Mode of TFEDREV, 2472.746674 539.000000 1.000000
 Spread - Variance, Standard deviation of TFEDREV, 150278906.593238 12258.829740
 Skew of TFEDREV, 32.430302
 Interquartile range of TFEDREV, 187.000000 539.000000 1649.000000

Descriptive Characteristics for TSTREV

Mean, Median, Mode of TSTREV, 13981.876206 4608.000000 40.000000
 Spread - Variance, Standard deviation of TSTREV, 2323006348.923632 48197.576173
 Skew of TSTREV, 19.203716
 Interquartile range of TSTREV, 1623.000000 4608.000000 11501.250000

Descriptive Characteristics for TLOCREV

Mean, Median, Mode of TLOCREV, 13800.836066 3524.000000 25.000000
 Spread - Variance, Standard deviation of TLOCREV, 2613455677.474381 51121.968638
 Skew of TLOCREV, 19.163061
 Interquartile range of TLOCREV, 1224.000000 3524.000000 10725.250000

Descriptive Characteristics for TOTALEXP

Mean, Median, Mode of TOTALEXP, 30560.999228 9660.000000 1638.000000
 Spread - Variance, Standard deviation of TOTALEXP, 11182506231.429939 105747.369856
 Skew of TOTALEXP, 19.065186
 Interquartile range of TOTALEXP, 3592.000000 9660.000000 25343.000000

Skew is far greater than 1, highlighting that the numbers for every column are skewed heavily towards right with long tail towards higher scale

► In []:

Plotting PMFS

► In [49]: *#Top 5 states with most total revenue*
findist_bystate.sort_values('TOTALREV',ascending=False)[0:5]

Out[49]:

	STATE	TOTALREV	ENROLL
32	NEW_YORK	27448356	1077381.0
4	CALIFORNIA	10329380	747009.0
13	ILLINOIS	5760419	437418.0
9	FLORIDA	3959408	375836.0
28	NEVADA	3220684	325990.0

► In [50]: *# Bottom 5 states with most total revenue*
findist_bystate.sort_values('TOTALREV',ascending=True)[0:5]

Out[50]:


	STATE	TOTALREV	ENROLL
45	VERMONT	84611	3992.0
26	MONTANA	117079	11348.0
19	MAINE	119828	8266.0
34	NORTH_DAKOTA	177525	12561.0
29	NEW_HAMPSHIRE	186461	17737.0

```
▶ In [51]: # Comparing 2 states -FLORIDA and VERMONT
flfindistdf = findistdf[findistdf.STATE == 'FLORIDA']
vfindistdf = findistdf[findistdf.STATE == 'VERMONT']
```


```
▶ In [52]: import thinkstats2
import thinkplot

flfindistdfpmf = thinkstats2.Pmf(flfindistdf['TOTALREV'], label='FLORIDA')
vfindistdfpmf = thinkstats2.Pmf(vfindistdf['TOTALREV'], label='VERMONT')
```


```
▶ In [53]: width=200000  
axis = [0, 800, 0, 0.0005]  
  
thinkplot.PrePlot(2, cols =2)  
thinkplot.Hist(flfindistdfpmf, align = 'right', width = width)  
thinkplot.Hist(vfindistdfpmf, align = 'left', width = width)  
thinkplot.Config(xlabel = 'Total Revenue', ylabel = 'PMF')
```



```
► In [54]: thinkplot.Pmf(flfindistdfpmf)  
thinkplot.Pmf(vfindistdfpmf)
```



```
► In [55]: thinkplot.PrePlot(2)
thinkplot.subplot(2)
#axis = [0, 800, 0, 0.0005]
thinkplot.Pmfs([flfindistdfpmf,vfindistdfpmf ])
thinkplot.Show(xlabel = 'Total Revenue', ylabel = 'PMF')
```



<Figure size 576x432 with 0 Axes>


Lets plot PMF of log transformed columns

```
► In [56]: findistdf.columns
```

```
Out[56]: Index(['STATE', 'ENROLL', 'NAME', 'YRDATA', 'TOTALREV', 'TFEDREV', 'TSTREV',
               'TLOCREV', 'TOTALEXP', 'lg_ENROLL', 'lg_TOTALREV', 'lg_TFEDREV',
               'lg_TSTREV', 'lg_TLOCREV', 'lg_TOTALEXP'],
              dtype='object')
```

```
▶ In [57]: lgflfindistdfpmf = thinkstats2.Pmf(flfindistdf['lg_TOTALREV'], label='FLORIDA')  
lgvfindistdfpmf = thinkstats2.Pmf(vfindistdf['lg_TOTALREV'], label='VERMONT')
```

```
▶ In [58]: width= 0.5  
#axis = [0, 800, 0, 0.0005]  
  
thinkplot.PrePlot(2, cols =2)  
thinkplot.Hist(lgflfindistdfpmf, align = 'right',color="blue", width = width)  
thinkplot.Hist(lgvfindistdfpmf, align = 'left',color="orange", width = width)  
thinkplot.Config(xlabel = 'Total Revenue (log)', ylabel = 'PMF')
```



```
▶ In [ ]:
```

```

In [59]: fltr4 = np.array(flfindistdf['lg_TOTALREV'].dropna())
vtr4 = np.array(vfindistdf['lg_TOTALREV'].dropna())

range_lb = int(min([np.min(fltr4), np.min(vtr4)]))
range_ub = int(max([np.max(fltr4), np.max(vtr4)]))

nbr_bins = range_ub - range_lb


pmf_fltr4 = np.histogram(np.array(fltr4),
                        bins=nbr_bins, range=(range_lb, range_ub))
pmf_vtr4 = np.histogram(np.array(vtr4),
                        bins=nbr_bins, range=(range_lb, range_ub))

width = 0.001
plt.bar(np.arange(range_lb, range_ub), pmf_fltr4[0], align = 'center', color="blue", label="FLORIDA")
plt.bar(np.arange(range_lb, range_ub) + width, pmf_vtr4[0], align = 'edge', color="orange", label="VERMONT")

plt.xlabel("Total Revenue (log)")
plt.ylabel("probability")
plt.legend(loc="best")

```

Out[59]: <matplotlib.legend.Legend at 0x225354ed668>



Based on the comparisons of PMF's Vermont - school districts are more likely to have lesser total revenues than Illinois school districts


In []:

Calculating CDF

```
In [60]: # CDF for total revenue
cdttlrev = np.sort(findistdf.TOTALREV) #ENROLL)
hist = np.histogram(cdttlrev, bins=15, range=(0,400000))
sz = len(cdttlrev)

plt.step(hist[1][:-1], np.cumsum(hist[0])/sz)
plt.xlabel("Total Revenue")
plt.ylabel("CDF(x)")
plt.title("CDF of TOTAL Revenue")
```

Out[60]: Text(0.5,1,'CDF of TOTAL Revenue')



95% of the total revenues for all school districts are less than 50,000

► In [61]: *# comparing CDF's of total revenue for school districts in different states*

```
vfindistdf = findistdf[findistdf.STATE == 'VERMONT']  
ilfindistdf = findistdf[findistdf.STATE == 'ILLINOIS']  
  
actvtr = np.array(vfindistdf['TOTALREV'].dropna())  
actiltr = np.array(ilfindistdf['TOTALREV'].dropna())
```

```


In [62]: # CDF for Vermont state
vcdttlrev = np.sort(activtr) #ENROLL)
histv = np.histogram(vcdttlrev, bins=20, range=(0,500000))
szv = len(vcdttlrev)

# CDF for Illionis state
ilcdttlrev = np.sort(actiltr) #ENROLL)
histil = np.histogram(ilcdttlrev, bins=20, range=(0,500000))
szil = len(ilcdttlrev)

plt.step(histv[1][:-1], np.cumsum(histv[0])/szv, color = 'blue', label = 'Total Revenue - Vermont')
plt.step(histil[1][:-1], np.cumsum(histil[0])/szil, color = 'red', label = 'Total Revenue - Illinois')
plt.xlabel("Total Revenue")
plt.ylabel("CDF(x)")
plt.title('CDF of Vermont Vs Illinois')
plt.legend()

```

Out[62]: <matplotlib.legend.Legend at 0x22539428400>



Overall school districts in Illinois have higher total revenue than Vermont and 98% of the total revenues for all school districts in Illinois is less than 100,000. Whereas for Vermont, almost 97% of the total revenues for school districts are below 50000\$. Put it in another way,

Illinois school districts have higher chance of having more Total Revenue.

► In []:



Plotting analytical distributions

```
► In [63]: # Plotting analytical distributions for total revenue

# calculate the mean and standard deviation

mean_tr = np.mean(findistdf.TOTALREV)
mean_tr

std_tr = np.std(findistdf.TOTALREV)
std_tr

# plot a normal distribution with the mean and standard deviation of total revenue


low = min(findistdf.TOTALREV)
high = max(findistdf.TOTALREV)

xs = np.linspace(low, high, 100)
ps = stats.norm.cdf(xs, mean_tr, std_tr)

plt.plot(xs, ps, label='model', color='0.6')

plt.title('Total Revenue vs CDF for normal distribution')
plt.xlabel('Sample of normal distribution')
plt.ylabel('CDF')
plt.legend()
```

Out[63]: <matplotlib.legend.Legend at 0x2253a063780>




```
► In [64]: # CDF for total revenue
cdttlrev = np.sort(findistdf.TOTALREV) #ENROLL)
hist = np.histogram(cdttlrev, bins=15, range=(0,max(findistdf.TOTALREV)))
sz = len(cdttlrev)

plt.plot(hist[1][:-1], np.cumsum(hist[0])/sz, color = 'b', label = 'smooth')

plt.step(hist[1][:-1], np.cumsum(hist[0])/sz, color = 'r', label = 'step')

plt.title('Total Revenue vs CDF for normal distribution')
plt.xlabel('Total Revenue')
plt.ylabel('CDF')
plt.legend()
```

Out[64]: <matplotlib.legend.Legend at 0x2253f19a550>




► In [65]: *# Overlapping distribution plots of normal distribution and cdf distribution of total revenues*

```
plt.plot(xs, ps, label='model', color='0.6')
plt.plot(hist[1][: -1], np.cumsum(hist[0])/sz, label='data')

plt.title('CDF of Total Revenue vs CDF of normal distribution')
plt.xlabel('Sample of normal distribution')
plt.ylabel('CDF')

plt.legend()
```

Out[65]: <matplotlib.legend.Legend at 0x2253f8cc5c0>



From the above overlapping plot of normal distribution vs total revenue, we can infer that Normal distribution doesn't represent the total revenue. Let us try to see if log normal distribution is applicable.

To find out I will use the log transformed total revenue column - `findistdf.lg_TOTALREV`


```
► In [66]: mean_ltr = np.mean(findistdf.lg_TOTALREV)
mean_ltr

std_ltr = np.std(findistdf.lg_TOTALREV)
std_ltr


# plot a normal distribution with the mean and standard deviation of total revenue

llow = min(findistdf.lg_TOTALREV)
lhigh = max(findistdf.lg_TOTALREV)

lxs = np.linspace(llow, lhigh, 10000)
lps = stats.norm.cdf(lxs, mean_ltr, std_ltr)

plt.plot(lxs,lps, label='model', color='0.6')
```

Out[66]: [



► In [67]:

```
mean_ltr = np.mean(findistdf.lg_TOTALREV)
mean_ltr

std_ltr = np.std(findistdf.lg_TOTALREV)
std_ltr


# plot a normal distribution with the mean and standard deviation of total revenue

llow2 = mean_ltr - 4 * std_ltr
lhigh2 = mean_ltr + 4 * std_ltr

lxs2 = np.linspace(llow2, lhigh2, 10000)
lps2 = stats.norm.cdf(lxs2, mean_ltr, std_ltr)

plt.plot(lxs2, lps2, label='model', color='0.6')
plt.title('Log (Total Revenue) vs CDF of log normal distribution')
plt.xlabel('Sample of log normal distribution')
plt.ylabel('CDF')
```

Out[67]: Text(0,0.5, 'CDF')




```
► In [68]: # CDF for log transformed total revenue
cdttlrev4 = np.sort(findistdf.lg_TOTALREV) #ENROLL)
hist = np.histogram(cdttlrev4, bins=1000, range=(0,max(findistdf.lg_TOTALREV)))
sz = len(cdttlrev4)

plt.plot(hist[1][:-1], np.cumsum(hist[0])/sz)

plt.title('Log (Total Revenue) vs CDF')
plt.xlabel('Log (Total Revenue)')
plt.ylabel('CDF')

#plt.legend()
```

Out[68]: Text(0,0.5, 'CDF')




► In [69]: *# Overlapping distribution plots of log normal distribution and cdf distribution of log (total revenues)*

```
plt.plot(lxs,lps, label='model', color='0.6')
plt.plot(hist[1][:,-1], np.cumsum(hist[0])/sz, label='data')

plt.title('CDF of Log(Total Revenue) vs CDF of log normal distribution')
plt.xlabel('Log(Total Revenue)')
plt.ylabel('CDF')
plt.legend()
```

Out[69]: <matplotlib.legend.Legend at 0x22541ee9ac8>

CDF of Log(Total Revenue) vs CDF of log normal distribution




From the above overlapping plot of log normal distribution vs log(total revenue), we can infer that log normal distribution perfectly fits for the variable total revenue

Probability plots for total revenue and Log (total revenue)

```
In [70]: xs = [-5, 5]
# y(x) = mean + std * x, here mean and standard deviation are from Total Revenue
ys = mean_tr + std_tr * np.sort(xs)
plt.plot(xs, ys, color='red', label='model')

plt.xlabel('Z')
plt.ylabel('Total Revenue (linear scale)')
```

Out[70]: Text(0,0.5,'Total Revenue (linear scale)')




```
In [71]: n = len(findistdf.TOTALREV)
xs2 = np.sort(np.random.normal(0, 1, n))
ys2 = np.sort(np.array(findistdf.TOTALREV))

plt.plot(xs2, ys2, color='blue', label='data')
plt.plot(xs, ys, color='red', label='model')

plt.title('Normal Probability plot on linear scale')
plt.xlabel("Z - standard deviation")
plt.ylabel('Total Revenue')
plt.legend()
```


Out[71]: <matplotlib.legend.Legend at 0x2253fb13fd0>



```
► In [72]: # Normal probability plot of log normal form

xs = [-5, 5]
# y(x) = mean + std * x, here mean and standard deviation are from log transformed Total Revenue
ys = mean_ltr + std_ltr * np.sort(xs)
plt.plot(xs, ys, color='red', label='model')
```

Out[72]: [<matplotlib.lines.Line2D at 0x2254151e588>]




```
In [73]: n = len(findistdf.lg_TOTALREV)
xs2 = np.sort(np.random.normal(0, 1, n))
ys2 = np.sort(np.array(findistdf.lg_TOTALREV))

plt.plot(xs2, ys2, color='blue', label='data')
plt.plot(xs, ys, color='red', label='model')

plt.title('Normal Probability plot on log scale')
plt.xlabel("Z - standard deviation")
plt.ylabel("Log (Total Revenue)")

plt.legend()
```

Out[73]: <matplotlib.legend.Legend at 0x22541859828>



From the above two normal probability plots, we can infer that data deviates substantially from normal model where as log normal model fits perfectly to the data with in 2 standard deviations (between -2 to 2) but deviates from the log normal model significantly for the school districts with lower and higher end of the (log) total revenue scale.

In []:


Scatter plots and Correlation analysis

In [74]: findistdf.head()

Out[74]:


	STATE	ENROLL	NAME	YRDATA	TOTALREV	TFEDREV	TSTREV	TLOCREV	TOTALEXP	lg_ENROLL	lg_TOTALREV	lg_TFED
0	ALABAMA	9609.0	AUTAUGA COUNTY SCHOOL DISTRICT	2016	80867	7447	53842	19578	76672	9.170455	11.300561	8.915
1	ALABAMA	30931.0	BALDWIN COUNTY SCHOOL DISTRICT	2016	338236	23710	145180	169346	299880	10.339514	12.731499	10.075
2	ALABAMA	912.0	BARBOUR COUNTY SCHOOL DISTRICT	2016	10116	2342	5434	2340	10070	6.815640	9.221874	7.755
3	ALABAMA	2842.0	EUFAULA CITY SCHOOL DISTRICT	2016	26182	3558	15900	6724	29843	7.952263	10.172827	8.175
4	ALABAMA	3322.0	BIBB COUNTY SCHOOL DISTRICT	2016	32486	3664	21846	6976	31662	8.108322	10.388565	8.205

```
In [75]: #plt.style.use('ggplot')
plt.style.use('seaborn')
findistdf.plot(x= 'TOTALREV', y = 'ENROLL', kind = 'scatter' )
plt.title('Scatter plot of Total Revenue Vs Enroll')
plt.show()
```




The above chart shows that school districts with higher total revenue has better enrollment than the school districts with lower total revenue, agrees with one of our assumptions.

```
► In [76]: # Scatter plot with jitter  
jitter = 20000  
TOTALREV = findistdf.TOTALREV + np.random.uniform(-jitter, jitter)  
plt.scatter(TOTALREV, findistdf.ENROLL, alpha = 0.3, s = 10)  
plt.title('Scatter plot of Total Revenue Vs Enroll')  
plt.show()
```



```
► In [77]: # Scatter plot between log(total revenue) and enroll

jitter = 1
lg_TOTALREV = findistdf.lg_TOTALREV + np.random.uniform(-jitter, jitter)
plt.scatter(lg_TOTALREV, findistdf.ENROLL, alpha = 0.3, s = 10)
plt.title('Scatter plot of log(Total Revenue) vs Enroll')
plt.show()
```



```
► In [78]: # Scatter plot between log(total revenue) and log(enroll)

jitter = 0.3
ENROLL4 = findistdf.lg_ENROLL + np.random.uniform(-jitter, jitter)
TOTALREV4 = findistdf.lg_TOTALREV + np.random.uniform(-jitter, jitter)
plt.scatter(TOTALREV4, ENROLL4, alpha = .2, s = 10)
plt.title('Scatter plot of log(Total Revenue) vs Enroll')
plt.show()
```



Characterizing Relationships

In [79]:


```
findistdf2 = findistdf.dropna(subset = ['ENROLL', 'TOTALREV'])  
bins = np.arange(0, 30000000, 200000)  
indicies = np.digitize(findistdf2.TOTALREV, bins)  
  
grps = findistdf2.groupby(indicies)
```

```
► In [80]: for i, group in grps:  
           print(i, len(group))
```

```
1 281561  
2 3932  
3 1093  
4 394  
5 220  
6 171  
7 87  
8 56  
9 41  
10 37  
11 31  
12 24  
13 18  
14 21  
15 17  
16 10  
17 3  
18 9  
19 3  
20 3  
21 2  
22 1  
24 1  
25 1  
26 2  
27 2  
28 2  
29 2
```

```
► In [81]: mean_tr = [group.TOTALREV.mean() for i, group in grps]  
           cdfs = [thinkstats2.Cdf(group.ENROLL) for i, group in grps]
```

```
In [82]: for percent in [75, 50, 25]:  
    enroll_percentiles = [cdf.Percentile(percent) for cdf in cdfs]  
    label = '%dth' % percent  
    thinkplot.Plot(mean_tr, enroll_percentiles, label=label)  
    thinkplot.Config(xlabel='Total Revenue', ylabel='ENroll', legend=True)  
    plt.title('Percentile plots of Total Revenue Vs Enroll')
```



Above percentiles plot of Total Revenue Vs Enroll, relationship is linear upto 3500000\$, after that relationship is going in the wrong direction.


```
► In [83]: descchar(findistdf['TOTALREV'])
```

```
Mean, Median, Mode of TOTALREV, 30255.458977 9715.000000 2276.000000  
Spread - Variance, Standard deviation of TOTALREV, 10677274625.544109 103330.898697  
Skew of TOTALREV, 18.316904  
Interquartile range of TOTALREV, 3653.000000 9715.000000 25200.250000
```

Covariance and Correlation

```
► In [84]: # Covariance  
  
def Cov(xs, ys, meanx=None, meany=None):  
    xs = np.asarray(xs)  
    ys = np.asarray(ys)  
  
    if meanx is None:  
        meanx = np.mean(xs)  
    if meany is None:  
        meany = np.mean(ys)  
  
    cov = np.dot(xs-meanx, ys-meany) / len(xs)  
    return cov  
  
Cov(findistdf2.TOTALREV, findistdf2.ENROLL)
```

```
Out[84]: 1024052236.613642
```

```
► In [85]: # Calculating Correlation between total revenue and Enroll in school districts
def Corr(xs, ys):
    xs = np.asarray(xs)
    ys = np.asarray(ys)

    meanx, varx = thinkstats2.MeanVar(xs)
    meany, vary = thinkstats2.MeanVar(ys)

    corr = Cov(xs, ys, meanx, meany) / np.sqrt(varx * vary)
    return corr

Corr(findistdf2.TOTALREV, findistdf2.ENROLL)
```

Out[85]: 0.9497119340376089

Correlation value of 0.95 indicates that total revenue and enroll variables are strongly and positively correlated; and it implies that school districts with higher total revenue tend to have higher enrollments in those schools. But our distributions are highly skewed and are not normal distributions, so let's find out the Spearman's Rank correlation as well.

```
► In [86]: def SpearmanCorr(xs, ys):
    xrank = pd.Series(xs).rank()
    yrank = pd.Series(ys).rank()
    return Corr(xrank, yrank)

SpearmanCorr(findistdf2.TOTALREV, findistdf2.ENROLL)
```

Out[86]: 0.9560835084463589

```
► In [ ]:
```

Hypothesis Testing

Defining Null Hypothesis: My earlier assumption is that school districts with higher total revenue will have higher enrollments in the school. Based on that, my Null hypothesis is that there is no relationship between Total revenue and school enrollments for school districts. The p-value for this correlation testing is to find out the probability of having such a high observed correlation of 0.95 by pure chance should be significant (pvalue > 0.05).

Let us find out with Hypothesis testing.

► In [87]: `findistdf.head()`

Out[87]:

	STATE	ENROLL	NAME	YRDATA	TOTALREV	TFEDREV	TSTREV	TLOCREV	TOTALEXP	lg_ENROLL	lg_TOTALREV	lg_TFED
0	ALABAMA	9609.0	AUTAUGA COUNTY SCHOOL DISTRICT	2016	80867	7447	53842	19578	76672	9.170455	11.300561	8.914
1	ALABAMA	30931.0	BALDWIN COUNTY SCHOOL DISTRICT	2016	338236	23710	145180	169346	299880	10.339514	12.731499	10.071
2	ALABAMA	912.0	BARBOUR COUNTY SCHOOL DISTRICT	2016	10116	2342	5434	2340	10070	6.815640	9.221874	7.754
3	ALABAMA	2842.0	EUFAULA CITY SCHOOL DISTRICT	2016	26182	3558	15900	6724	29843	7.952263	10.172827	8.174
4	ALABAMA	3322.0	BIBB COUNTY SCHOOL DISTRICT	2016	32486	3664	21846	6976	31662	8.108322	10.388565	8.201

```
► In [88]: findist_bystate_mn = pd.DataFrame()
findist_bystate_mn = findistdf.groupby('STATE', as_index = False)['TOTALREV', 'ENROLL'].mean()
findist_bystate_mn.sort_values('TOTALREV', ascending=False)[0:5]
```

Out[88]:

	STATE	TOTALREV	ENROLL
9	HAWAII	2.047635e+06	183417.458333
18	MARYLAND	4.132408e+05	35081.586806
7	FLORIDA	3.172400e+05	37111.053483
26	NEVADA	1.803720e+05	21480.924020
30	NORTH_CAROLINA	9.168736e+04	11389.473421

```
► In [89]: findist_bystate_mn.sort_values('ENROLL', ascending=False)[0:5]
```

Out[89]:

	STATE	TOTALREV	ENROLL
9	HAWAII	2.047635e+06	183417.458333
7	FLORIDA	3.172400e+05	37111.053483
18	MARYLAND	4.132408e+05	35081.586806
26	NEVADA	1.803720e+05	21480.924020
41	UTAH	8.064527e+04	12644.656670

```
► In [90]: Hyp_df = findistdf.dropna(subset = ['ENROLL', 'TOTALREV'])  
Hyp_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 287744 entries, 0 to 358292  
Data columns (total 15 columns):  
STATE          287744 non-null object  
ENROLL         287744 non-null float64  
NAME           287744 non-null object  
YRDATA         287744 non-null int64  
TOTALREV       287744 non-null int64  
TFEDREV        287744 non-null int64  
TSTREV         287744 non-null int64  
TLOCREV        287744 non-null int64  
TOTALEXP       287744 non-null int64  
lg_ENROLL      287744 non-null float64  
lg_TOTALREV    287744 non-null float64  
lg_TFEDREV     287744 non-null float64  
lg_TSTREV      287744 non-null float64  
lg_TLOCREV     287744 non-null float64  
lg_TOTALEXP    287744 non-null float64  
dtypes: float64(7), int64(6), object(2)  
memory usage: 35.1+ MB
```

```
► In [91]: # Pearson correlation  
pecorr, p = stats.pearsonr(Hyp_df.TOTALREV, Hyp_df.ENROLL)  
pecorr, p
```

Out[91]: (0.9497119340376095, 0.0)

```
► In [92]: # Spearman correlation  
corr2, p2 = stats.spearmanr(Hyp_df.TOTALREV, Hyp_df.ENROLL)  
corr2, p2
```

Out[92]: (0.9560835084463587, 0.0)

```
▶ In [93]: trenr_df = Hyp_df[['TOTALREV', 'ENROLL']]
```

```
▶ In [94]: stat, p, dof, expected = stats.chi2_contingency(trenr_df)
          stat, p, dof, expected
```

```
Out[94]: (108740524.95546971, 0.0, 287743, array([[8.18934126e+04, 8.58258740e+03],
        [3.34147679e+05, 3.50193205e+04],
        [9.98187977e+03, 1.04612023e+03],
        ...,
        [1.23614918e+04, 1.29550816e+03],
        [6.90078449e+03, 7.23215508e+02],
        [2.81317395e+03, 2.94826049e+02]]))
```

Method for testing this null hypothesis is to randomly generate values for total revenue and enroll with the same mean and standard deviation of the current data set and calculate the Correlation and P- value for that sample data set. Repeat the process for some iterations (100)

In [95]:

```
def samplepermute(iters = 100):  
    '''  
    Function to permute the TOTALREV randomly and calculate Correlation, p value and Covariance of that data  
    iters is number of iterations of test  
    returns:  
    smpcorr - sample correlation  
    smppval - sample pvalue  
    count/iters - % of samples that have sampled correlation greater than observed correlation  
    '''  
    smpcorr = []  
    smppval = []  
    count = 0  
    df = pd.DataFrame()  
    for j in range(iters):  
  
        corr, p = stats.pearsonr(np.random.permutation(Hyp_df.TOTALREV),Hyp_df.ENROLL)  
        cov = Cov(np.random.permutation(Hyp_df.TOTALREV), Hyp_df.ENROLL)  
  
        if abs(corr) >= pecorr:  
            count += 1  
            smpcorr.append(corr)  
            smppval.append(p)  
    return smpcorr, smppval, count/iters
```

In [96]:

```
test2corr, test2p, test2count = samplepermute( iters = 100)  
test2count
```

Out[96]: 0.0

The probability of having such a high (observed - 0.95) correlation between Total Revenue & Enroll by chance is 0. Hence null hypothesis that there is no correlation between Total Revenue and Enroll is false.

► In []:


Regression Analysis

► In [97]:

```
findistdf1 = findistdf.dropna(subset = ['ENROLL', 'TOTALREV'])
```




```
► In [98]: #Plotting scatter plot between Total Revenue and Enroll  
#plt.style.use('ggplot')  
plt.style.use('seaborn')  
findistdf1.plot(x= 'TOTALREV', y = 'ENROLL', kind = 'scatter' )  
plt.xlabel('Total Revenue')  
plt.ylabel('Enroll')  
plt.title('Scatter plot of Total Revenue Vs Enroll with Best fit line')  
plt.show()
```



Linear Least Square Model

```
► In [99]: # plotting scatter plot again
plt.style.use('seaborn')
findistdf1.plot(x= 'TOTALREV', y = 'ENROLL', kind = 'scatter' )
plt.xlabel('Total Revenue')
plt.ylabel('Enroll')
plt.title('Scatter plot of Total Revenue Vs Enroll')
plt.show()
```



► In [100]: *#Calculating slope & iter*

```
meantr = np.mean(findistdf1.TOTALREV)
meanenr = np.mean(findistdf1.ENROLL)

slope = (meantr * meanenr - np.mean(findistdf1.TOTALREV * findistdf1.ENROLL)) / (meantr**2 - np.mean(findistdf1.TOTALREV**2))

inter = meanenr - meantr*slope

meantr, meanenr, slope, inter
```

Out[100]: (30255.458977424376, 3170.830383952402, 0.09590984885567204, 269.0338863686443)


► In [101]: *# Drawing best fit line*

```
plt.scatter(findistdf1.TOTALREV, findistdf1.ENROLL)
TR_minmax = [np.min(findistdf1.TOTALREV), np.max(findistdf1.TOTALREV)]

Regressionline = [slope*t + inter for t in TR_minmax ]

plt.plot(TR_minmax, Regressionline, color = "red", alpha =0.3)
plt.xlabel('Total Revenue (Independent variable)')
plt.ylabel('Enroll (dependant variable)')
plt.title('Scatter plot of Total Revenue Vs Enroll with Linear line fit')
```

Out[101]: Text(0.5,1,'Scatter plot of Total Revenue Vs Enroll with Linear line fit')



Above plot confirms the linear relationship between Total Revenue and Enroll

Goodness of Linear Least Square fit

Calculating Residuals, RMSE, Coefficient of Determination

Goodness of linear least square fit can be found by comparing the Root mean square error between with model and without model.

```
▶ In [102]: # Plotting residuals
def Residuals(xs, ys, inter, slope):
    xs = np.asarray(xs)
    ys = np.asarray(ys)
    res = ys - (inter + slope * xs)
    return res

findistdf1['residual'] = Residuals(findistdf1.TOTALREV, findistdf1.ENROLL, inter, slope)
#len(findistdf1['residual'])
```

```
▶ In [103]: # Function to calculate Root mean squared error
def RMSE(ys, pred_ys):
    yactual = np.array(ys)
    ypred = np.array(pred_ys)
    error = (yactual - ypred)**2
    errmean = np.mean(error)
    errsqr = sqrt(errmean)
    return errsqr

xs = np.array(findistdf1.TOTALREV)
# predicted value for enroll using linear least square
pred_ys = [inter + (slope * x) for x in np.array(findistdf1.TOTALREV)]
RMSE(findistdf1.ENROLL, pred_ys)
```

Out[103]: 3267.521506975128

```
► In [104]: # Calculating Root Mean Square Error (RMSE) - Standard deviation of residuals
np.std(findistdf1['residual']), np.std(findistdf1['ENROLL'])
```

Out[104]: (3267.521506975102, 10435.199664236075)

Without any model, RMSE of predicted Enroll numbers is represented by its standard deviation – which here in this case is 10435.

With Linear Least Square fit model, RMSE of predicted Enroll numbers from known Total Revenues is calculated by finding the residuals from prediction (Observed Enroll – Predicted Enroll) and finding the standard deviation from the residual. In this case it is 3267.

As predicting Enrollment numbers with Linear Least Square model results in lesser standard deviation, in this case knowing the total revenue and predicting enrollment numbers from it has significantly helped for better prediction and reducing the error.

```
► In [105]: # Coefficient of Determination:
resid_var = np.var(findistdf1['residual'])
enroll_var = np.var(findistdf1['ENROLL'])

CoeffD = 1 - resid_var/enroll_var
CoeffD
```

Out[105]: 0.9019527576534582

CoeffD of 0.90 indicates that total revenue helps predict almost 90% of the variance in the enrollment numbers for school districts.

Plotting residuals

```
▶ In [106]: bins = np.arange(min(findistdf1.TOTALREV), max(findistdf1.TOTALREV), 250000)
indices = np.digitize(findistdf1.TOTALREV, bins)
groups = findistdf1.groupby(indices)


trbin_means = [group.TOTALREV.mean() for _, group in groups][1:-1]
#len(trbin_means) = 50
```

```
▶ In [107]: cdfs = [thinkstats2.Cdf(group.residual) for _, group in groups][1:-1]
```

```
▶ In [108]: def PlotPercentiles(trbin_means, cdfs):
    thinkplot.PrePlot(3)
    for percent in [75, 50, 25]:
        weight_percentiles = [cdf.Percentile(percent) for cdf in cdfs]
        label = '%dth' % percent
        thinkplot.Plot(trbin_means, weight_percentiles, label=label)
```

```
► In [109]: PlotPercentiles(trbin_means, cdfs)

thinkplot.Config(xlabel="Total Revenue", ylabel='Residuals')
```




The residual plots are not straight lines, indicates that relationship between total revenue and enroll is non - linear. The gap between inter quartile residuals is most at the total revenue of 3 Million.

```
► In [ ]:
```



```
► In [110]: # Plotting best fit line with stats modules
plt.plot(findistdf1.TOTALREV, findistdf1.ENROLL, 'o', label='original data')
plt.plot(findistdf1.TOTALREV, inter + slope*findistdf1.TOTALREV, 'r', label='fitted line')
plt.legend()
plt.show()
```



```
► In [ ]:
```

Testing Linear Model

► In [111]: *#To estimate the sampling distribution of inter and slope, I'll use resampling.*

```
def SampleRows(df, nrows, replace=False):  
    """Choose a sample of rows from a DataFrame.  
  
    df: DataFrame  
    nrows: number of rows  
    replace: whether to sample with replacement  
  
    returns: DataDf  
    """  
    indices = np.random.choice(df.index, nrows, replace=replace)  
    sample = df.loc[indices]  
    return sample  
  
def ResampleRows(df):  
    """Resamples rows from a DataFrame.  
  
    df: DataFrame  
  
    returns: DataFrame  
    """  
    return SampleRows(df, len(df), replace=True)
```

```
▶ In [112]: def SamplingDistributions(findistdf1, iters=101):  
    inters = []  
    slopes = []  
    for _ in range(iters):  
        sample = ResampleRows(findistdf1)  
        TOTALREV = sample.TOTALREV  
        ENROLL = sample.ENROLL  
  
        slope = (meantr * meanenr - np.mean(findistdf1.TOTALREV * findistdf1.ENROLL)) / (meantr**2 - np.mean(f  
        slopes.append(slope)  
        inter = meanenr - meantr*slope  
        inters.append(inter)  
  
    return inters, slopes
```

```
▶ In [113]: inters, slopes = SamplingDistributions(findistdf1, iters=100)
```


```
▶ In [114]: slope_cdf = thinkstats2.Cdf(slopes)  
pvalue = slope_cdf[0]  
pvalue
```

```
Out[114]: 0.0
```

Probability that the slope in the sampling distribution falls below 0 (p-value) is 0; as it is less than 0.001 indicating that the relation between Total Revenue and Enroll is statistically significant and not by chance.

```
► In [115]: for slope, inter in zip(slopes, inters):  
            fxs = np.sort(findistdf1.TOTALREV)  
            fys = inter + slope * fxs  
            thinkplot.Plot(fxs, fys, color='gray', alpha=0.01)  
  
            thinkplot.Config(xlabel="Total Revenue", ylabel='Residuals', title = "Regression lines for various iterations")
```

No handles with labels found to put in legend.



After repeated sampling the regression line roughly stayed in the same place, so it is a low variance model

```
► In [ ]:
```

Regression Analysis - Ordinary Least Square Model


```

In [116]: import statsmodels.formula.api as smf

formula = 'findistdf1.ENROLL ~ findistdf1.TOTALREV'
model = smf.ols(formula, data=findistdf1)
results = model.fit()
results.summary()

```

Out[116]: OLS Regression Results

Dep. Variable:	findistdf1.ENROLL	R-squared:	0.902
Model:	OLS	Adj. R-squared:	0.902
Method:	Least Squares	F-statistic:	2.647e+06
Date:	Sun, 03 Mar 2019	Prob (F-statistic):	0.00
Time:	16:59:57	Log-Likelihood:	-2.7367e+06
No. Observations:	287744	AIC:	5.473e+06
Df Residuals:	287742	BIC:	5.473e+06
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	269.0339	6.347	42.387	0.000	256.594	281.474
findistdf1.TOTALREV	0.0959	5.9e-05	1626.956	0.000	0.096	0.096

Omnibus:	167866.320	Durbin-Watson:	1.533
Prob(Omnibus):	0.000	Jarque-Bera (JB):	2577381320.624
Skew:	0.853	Prob(JB):	0.00
Kurtosis:	466.649	Cond. No.	1.12e+05

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.12×10^5 . This might indicate that there are strong multicollinearity or other numerical problems.

```
▶ In [117]: inter = results.params['Intercept']  
            slope = results.params['findistdf1.TOTALREV']  
            inter, slope
```

Out[117]: (269.0338863686441, 0.09590984885567201)

Interpreting the coefficients: slope value of 0.095 infers that unit increase in total revenue is associated with 0.095 unit increase in enroll numbers for the school districts.

```
▶ In [118]: slope_pvalue = results.pvalues['findistdf1.TOTALREV']  
            slope_pvalue
```

Out[118]: 0.0

As P-value is less than 0.001, the estimated slope is significant

```
▶ In [119]: # Coefficient of determination  
            results.rsquared
```

Out[119]: 0.9019527576534563


R- Square value of 0.90 shows that variation in enrollment can be explained by variation in Total Revenue upto 90%. As more variance is being explained by the model, once again proves the fit of the model.

► In [120]: *#plotting residuals*

```
pred_val = results.fittedvalues.copy()
true_val = findistdf1['ENROLL'].values.copy()
residual = true_val - pred_val
```



```
► In [121]: fig, ax = plt.subplots(figsize=(10,10))  
residplot = ax.scatter(residual, pred_val)
```




```
In [122]: # Drawing normal probability plot

fig, ax = plt.subplots(figsize=(8,6))
_, (_, __, r) = stats.probplot(residual, plot=ax, fit=True)

r**2
#stats.probplot(residual, plot=ax, fit=True)
```

Out[122]: 0.28856453419639605



Above normality plot indicates that this model is good fit only between quartile -2 to +2. After that it significantly deviates from linear model.

```
► In [123]: # Confidence intervals
           results.conf_int()
```

Out[123]:

	0	1
Intercept	256.593660	281.474113
findistdf1.TOTALREV	0.095794	0.096025

```
► In [124]: # p-values for the model coefficients
           results.pvalues
```

Out[124]: Intercept 0.0
findistdf1.TOTALREV 0.0
dtype: float64

Again p-values are way less than 0.05, indicating that the relation between dependent and independent variables is genuine.

Correlation matrix


```
► In [125]: # Compute the correlation matrix
           corr = findistdf1.corr()
```

```
► In [126]: # Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
```

```
Out[126]: <matplotlib.axes._subplots.AxesSubplot at 0x225357cbcf8>
```



▶ In []:

Multiple Linear Regression

```

In [127]: # understanding whether the enrollment in school districts depends on Total Revenue and STATE as well
formula2 = 'findistdf1.ENROLL ~ findistdf1.TOTALREV + findistdf1.STATE'
model2 = smf.ols(formula2, data=findistdf1)
results2 = model2.fit()
results2.summary()

```

Out[127]: OLS Regression Results

Dep. Variable:	findistdf1.ENROLL	R-squared:	0.909
Model:	OLS	Adj. R-squared:	0.909
Method:	Least Squares	F-statistic:	6.001e+04
Date:	Sun, 03 Mar 2019	Prob (F-statistic):	0.00
Time:	17:00:19	Log-Likelihood:	-2.7256e+06
No. Observations:	287744	AIC:	5.451e+06
Df Residuals:	287695	BIC:	5.452e+06
Df Model:	48		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1532.3320	56.355	27.191	0.000	1421.879	1642.785

In []:

Comparing model 1 vs model2:

1) Adjusted R-square of model2 is .909, better than Adjusted R-square of model1 (0.902) - indicating that model2 can explain slightly more variation in dependent variable compared to model1.

2) AIC of model1 is 5.473e+06, which is slightly higher than AIC of model2 5.451e+06, indicating that model2 (enrollment as a function of total revenue and state) is slightly better model among the 2 models.

- 3) HAWAII has the highest absolute coefficient - indicating that ENROLLment numbers change hugely with a single unit of variation in TOTALREV for that state. In other words, we can probably see more enrollment numbers for every same number of units increase in total revenue compared to all other states (all remaining things being constant).
- 4) LOUISIANA has the lowest absolute coefficient value - indicating that ENROLL numbers will change at a slower rate compared to all other states for the same unit of increase in the total revenue (all remaining things being constant).
- 5) Overall adding STATE to the ordinary least square model, improved the model very slightly but not significantly. But interesting aspect of adding STATE to the equation is it gives us insights into how each is the relationship between Total Revenue and Enroll for each STATE.

► In []:

Outcome of EDA:

The assumption that I had before exploring this data set was that the school districts that are in higher revenue states will have more chance of higher enrollments in the school. After performing EDA, I did find statistical correlation between Total Revenues and Enrollment of the school districts. So my assumption was correct. One more observation that I made is not all states will respond similarly to the total revenue numbers. For example, take state like LOUISIANA, even though Total Revenues increase for this state, Enrollment numbers will not raise proportionally when compared to other states.

What was missing during the analysis?

Having demographic information of each school districts like population, family size, number of school going kids, family income etc. would have added more sense to the analysis. Also one of these or couple of these could be the confounding factors that I highlighted above with STATE – LOUISIANA.

Variables that could have helped in the analysis?

As states above, demographic information could have helped more in the analysis in finding the actual enrollment prediction for the school districts.

Assumptions made correct or incorrect?

No. The assumptions that I made that enrollments in school districts are based on Total Revenues of the state was correct, backed by the higher correlation factor and the linear models

Challenges faced

In the selected data, there were some other aggregated data sets available, which I wanted to explore and compare with the financial school district data set that I selected. But some of the variables that I wanted to explore like GRADES_ALL_G etc. are not available through those sheets. I feel enrollment numbers depend a lot on population or number of families living in that school district. Having a demographic data for these school district would have been an interesting analysis, I would like to try. But that information was not available readily, so I couldn't venture into that analysis. One more challenge is that some of the states like California, New York are too big to be compared with other smaller states like Vermont, but before realizing this when I performed EDA overall financial school district data set, numbers were highly skewed – and didn't find an ideal number of bins to represent these variables into proper histogram. Plotting PDF and CDF without using the thinkstats2, thinkplot modules has been a challenge, at some places I gave up trying to figure out other means and ended up using these modules. I would like to explore these in free time.

► In []:

